

logistic-regression-model

December 29, 2023

```
[36]: import numpy as np
      from matplotlib.pyplot import plot
      import pandas
      from sklearn.model_selection import train_test_split
```

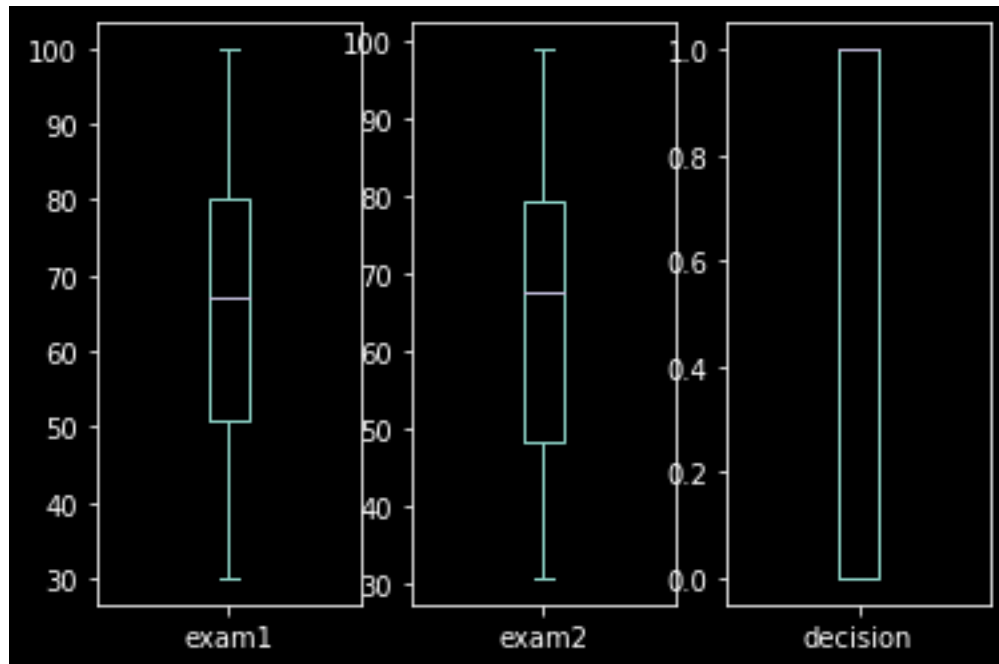
```
[37]: # Load data from file

      data = pandas.read_csv('acceptance_data.txt', names = [
          ↪['exam1', 'exam2', 'decision'])

      print(data.describe())
      data.plot(kind='box', subplots=True)
```

	exam1	exam2	decision
count	100.000000	100.000000	100.000000
mean	65.644274	66.221998	0.600000
std	19.458222	18.582783	0.492366
min	30.058822	30.603263	0.000000
25%	50.919511	48.179205	0.000000
50%	67.032988	67.682381	1.000000
75%	80.212529	79.360605	1.000000
max	99.827858	98.869436	1.000000

```
[37]: exam1      AxesSubplot(0.125,0.125;0.227941x0.755)
      exam2      AxesSubplot(0.398529,0.125;0.227941x0.755)
      decision  AxesSubplot(0.672059,0.125;0.227941x0.755)
      dtype: object
```



```
[38]: X = data.drop(columns=['decision']).round(2) # Features (exam scores)
      y = data['decision'] # Labels (acceptance decision)
```


```
[39]: # Perform data preprocessing and split data to 2 sets

      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25,
      ↪random_state=100)

      print(type(X_train))
      X_train.plot(kind='box', subplots=True, sharey=True)
```

```
<class 'pandas.core.frame.DataFrame'>
```


```
[39]: exam1      AxesSubplot(0.125,0.125;0.352273x0.755)
      exam2      AxesSubplot(0.547727,0.125;0.352273x0.755)
      dtype: object
```



```
[40]: # Feature scaling (Min-Max Scaling)
xmin = X_train.min()
xmax = X_train.max()
X_train = (X_train - xmin) / (xmax - xmin)
X_test = (X_test - xmin) / (xmax - xmin)

# Box plot after feature scaling
X_train.plot(kind='box', subplots=True, sharey=True)
```

```
[40]: exam1      AxesSubplot(0.125,0.125;0.352273x0.755)
exam2      AxesSubplot(0.547727,0.125;0.352273x0.755)
dtype: object
```



```
[41]: # Convert DataFrame to NumPy arrays
```

```
X_train = X_train.to_numpy()
X_test = X_test.to_numpy()
y_train = y_train.to_numpy()
y_test = y_test.to_numpy()
```

```
[42]: # Sigmoid function for logistic regression
```

```
def sigmoid(z):
    return 1 / (1 + np.exp(-z))
```

```
# Cost function for logistic regression
```

```
def cost_function(theta, X, y):
    m = len(y)
    h = sigmoid(X @ theta)
    cost = -(1/m) * np.sum(y * np.log(h) + (1 - y) * np.log(1 - h))
    return cost
```

```
# Gradient descent for parameter optimization
```

```
def gradient_descent(theta, X, y, alpha, iterations):
    m = len(y)
    for _ in range(iterations):
        h = sigmoid(X @ theta)
        gradient = (1/m) * X.T @ (h - y)
        theta -= alpha * gradient
    return theta
```

```
[43]: # Initialize parameters
theta_initial = np.zeros(X_train.shape[1])

# Set hyperparameters
learning_rate = 0.01
num_iterations = 1000
```

```
[44]: # Train the logistic regression model
theta_optimized = gradient_descent(theta_initial, X_train, y_train,
    ↪ learning_rate, num_iterations)
```

```
[45]: # Assess the fitted model on the test data

# Prediction function
def predict(theta, X):
    return (sigmoid(X @ theta) >= 0.5).astype(int)

# Make predictions on the test set
y_pred = predict(theta_optimized, X_test)
```

```
[46]: # Calculate accuracy
accuracy_v1 = np.mean(y_pred == y_test)
print(f"Accuracy on test set: {accuracy_v1*100}%")
```

Accuracy on test set: 56.00000000000001%

```
[48]: # Step 6: Generate predictions for new data
# (Assuming you have new data in a variable 'new_data')
# Preprocess new_data similar to the training data
# Add bias term
# Use the trained theta_optimized to make predictions
```